

Social Trust

Phil 101 · Day 9

Brian Weatherson

2026-09-29

Where we are

Recap

Anderson suggested that we have a way of making sensible decisions about what to believe in fields we can't really follow: make judgments about the experts.

- ▶ Are they **in a position to know**?
- ▶ Are they **honest**?
- ▶ Are they **responsive** to people who disagree with them?

When the trustworthy experts converge, that's good enough.

Follow up questions

First: *why should a consensus of experts be worth anything in the first place?* Anderson tells you to look for one. She doesn't say what makes it reliable.

Second: what happens when the machinery that generates your trust gets turned against you?

These are related; one reason for being sceptical of a consensus of experts is that you think the machinery has already been turned against you.

A quick iClicker reading question

From the Nguyen reading. Which is harder to get out of?

- A. An epistemic bubble
- B. An echo chamber
- C. They're two names for the same thing

Why trust science

Middle school science version

Most people, asked why science is trustworthy, reach for the same thing: **the Scientific Method**. Not surprising; that's literally on the middle school syllabus in many places (including AA).

You know the basic picture: Observe, hypothesize, test, conclude. That gets you from data to theories, and it's all great.

Most people who look in detail at how modern science works are very sceptical that there is such a clean story.

It isn't one thing

We looked at a small piece with Naomi Oreskes, who wrote a prominent book on modern science. Here's what she says about scientific method.

i Oreskes

Essentially my argument is: It isn't one thing.

Different sciences work differently. Geology doesn't run experiments the way chemistry does. Epidemiology can't randomize most of what it studies.

If there is one golden **Scientific Method** then we'd have to start asking which sciences follow it, because it can't be all of them, and then just trust those.

Oreskes' alternative

What the sciences have in common is a **social arrangement**.

A claim gets handed to a community of people trained to attack it, who have professional incentives to succeed in those attacks, and only what survives that gets kept.

You don't trust any one scientist. You trust that the people who tried to knock something down were willing and able to find any mistakes, and their non-discovery is good evidence.

Five pillars

i Oreskes

I didn't start out thinking there would be five key elements, but that's where I got to: consensus, diversity, method, evidence, and values.

Method here does not mean the One Scientific Method, and it isn't the whole story.

Diversity

This wasn't as much of a political buzz-word when Oreskes was writing 8-10 years ago! But you might worry about it for a couple of reasons.

1. Requiring diversity, even of opinions, sounds like politics not science.
2. The aim is **consensus**; encouraging diverse opinions seems like it makes that harder.

The point is that we should avoid blind spots. We shouldn't all be testing theories the same way, or for the same flaw.

Criticism only works if the critics are different enough from each other to catch different errors. Homogeneity is **ineffective**.

Values

i Oreskes

All people have values, and we always will have values... if you had scientists with no values, that would be truly scary.

The demand that science be value-free is not achievable, and Oreskes thinks it isn't even desirable. What you want is that the values are visible and open to challenge, like everything else.

Consensus

So here's why consensus is valuable. Theories that are the consensus view have run the gauntlet of criticisms, and lived to tell the tale.

If that's right, and you should never take anything like this at face value but I think it's not a terrible picture of how modern science works, it has two consequences.

1. We now have an affirmative reason to believe in the consensus; it survived tough criticism.
2. We have some pointers to when a consensus, or non-consensus, is **fake**.

Trust

Two bad outcomes

C. Thi Nguyen thinks we use one phrase for two very different problems, and that the confusion makes us reach for a fix that can't work.

Epistemic bubble

An informational network from which relevant voices have been **excluded by omission**.

Nobody did anything to you. You follow the people you follow, the feed shows what it shows, and some views simply never arrive.

Echo chamber

A social structure from which other relevant voices have been **actively discredited**.

Here somebody *has* done something. You've been brought to distrust the outside, so its experts turn up already labelled as liars or fools.

Why the difference matters

The escape strategies are different in the two cases.

Bubbles are a gap, so they can be remedied, filled in, by simple exposure.

i Nguyen

Epistemic bubbles are rather ramshackle; they go up easily, and they collapse easily, too.
Echo chambers are far more pernicious and far more robust.

Escape from the chamber

Here's the thing that makes chambers hard to get out of. Members have prepped responses to the other side.

i Nguyen

They read — but do not accept — mainstream and liberal news sources. They hear, but dismiss, outside voices.

So the standard remedy, get people out of their bubble, show them the other view, has been pre-emptively rejected.

Pre-emption

The mechanism has a name. Endre Begby calls it **evidential pre-emption**.

You're told in advance what the other side will say, and *why* they'll say it. So when they say it, they're doing exactly what you were warned they would do.

Nguyen illustrates this with a radio host from Past Times. Here's an invented case, closer to home.

Two channels

Both of these are made up, and I'm not endorsing either.

Overclock is a streamer. They say LLMs are the biggest thing since the printing press, and the people telling you otherwise are protecting their own jobs.

Slow Down is a podcast. They say the big AI companies are a bubble built on other people's work, and the people telling you otherwise are paid to.

The warning comes first

Note Each of them tells you, *before* you ever meet the other side, what the other side is going to say and why.

- ▶ Overclock: *You'll hear professors say it can't really reason. They've said that about every tool that ever threatened them.*
- ▶ Slow Down: *You'll hear researchers say it's nearly there. Every one of them holds equity.*

Now go and read the professor who is critical of AI, and go and read the researcher who is enthusiastic.

The trap closes

The professor says it can't really reason. Which is what you were told she would say.

The researcher says it's nearly there. Which is what you were told he would say.

Whatever the outsider says, the insider called it first.

When you are presented with the refutation from the 'other' side, your first reaction is that it is *confirmation* of what your side said about the other.

Why it works

Notice that neither warning is a lie.

Some academics really are threatened.¹ Some researchers really do hold equity.² The warnings are plausible, and that is why they work.

The trick isn't inventing a grand conspiracy.

It's planning for what the other side will say, and coming up with sweeping comments, typically about the **character** of the other side, that work no matter what.

¹Not me; I've never met a chatbot that can replicate my typos.

²Also not me; my money is in boring index funds.

Bubbles and chambers

If you'd simply never come across critical takes, you'd be in a **bubble**. Someone sends you a good take, you read it, you think again. Done.

If you've been told the critics are bought, you're in a **chamber**. Someone sends you the same good take, you check the author, you find she has a grant from a tech company, or works in a field that's losing jobs to AI, and you've got a reason to dismiss their work.

Getting out

If the trouble is who you trust rather than what you've seen, no single piece of evidence can fix it.

Nguyen's answer is a **social-epistemic reboot**: suspend the whole network of trust and rebuild it, rather than patching one belief.¹

i Nguyen

Black left the movement by, basically, performing a social reboot. He completely abandoned everything he'd believed in, and spent years building a new belief system from scratch.

¹This was Descartes' dramatic solution for the problem that our senses sometimes deceive us. It's a little more plausible here.

Why chambers work

A chamber issues its own credentials and stages its own consensus.

It teaches you to read a critic's honesty, any time they report on their background or raise worries about some detail, as proof of bad faith.

Those are Anderson's most important criteria. The chamber doesn't try to stop you assessing trustworthiness. It tries to capture the assessment.

Capturing vigilance

Sperber's answer to *aren't we all dupes* was that we run checks as we listen. Vigilance is what keeps you from being taken in.

i Nguyen

Once you're in, the echo chamber's belief systems function as a trap, making future acts of intellectual vigilance only reinforce the echo chamber's worldview.

The thing that was meant to protect you is now making things worse. This isn't an easy problem!

iClicker Poll

Pick something you're confident about. Could you say what would change your mind?

- A. Yes, and I could name the specific evidence
- B. Yes, but only vaguely
- C. No, not really

Module Quiz 1

Opens after class, on Lectures 1 to 9, due Sunday. Open book, multiple choice.

It's worth 2.5%, so it's really more important as a check on how you're doing rather than a big part of the assessment.

Thursday

We've spent two weeks on knowledge that comes from other people, and it's time to return to the thing that's normally central to philosophical discussions of knowledge: perception.

We're doing it last because it's a good segue into the second unit of the course, on philosophical views about minds, and the long running debate about the relationship between mind and body.